

Initial image + task

Pretrained q_ψ


u_1 : Grasp at τ_1

u_2 : Lift at τ_2

Learned visual transition model q_ϕ

Shared across steps

Predicted states + times




x_0

Initial state

q_ϕ

$u_1, \Delta\tau_1$




$\hat{x}_1$

Attach, τ_1

q_ϕ


$u_2, \Delta\tau_2$



$\hat{x}_2$

Lift, τ_2


Local event + elapsed time



Initial image + task

Timed state memory

Video dynamics q_θ



All frames generated jointly
Soft state conditions